

## Pandas vs. Numpy

What if we want to inspect the data on Abraham Lincoln in 'height\_age\_arr' but cannot remember his integer position. Is there a convenient way to access the data by indexing the name of the president like:

```
print(height_age_arr['Abraham Lincoln'])
```

PY

Unfortunately, we will receive an error message. However, it is possible to do this in pandas. The pandas library is built on top of numpy, meaning a lot of features, methods, and functions are shared.

By convention, import the library under a short name "pd":

```
import pandas as pd
```

PY

In later lessons, we will see that pandas allows us to access data by indexing the name directly.



As numpy ndarrays are homogeneous, pandas relaxes this requirement and allows for various dtypes in its data structures.

## Series

The Series is one building block in pandas. **Pandas Series** is a one-dimensional labeled array that can hold data of any type (integer, string, float, python objects, etc.), similar to a column in an excel spreadsheet. The axis labels are collectively called **index**.

If we are given a bag of letters a, b, and c, and count how many of each we have, we find that there are 1 a, 2 b's, and 3 c's. We could create a Series by supplying a list of counts and their corresponding labels:

```
pd.Series([1, 2, 3], index=['a', 'b', 'c']) # with index
```

PY

Try it Yourself

Alternatively, the values can be a numpy array:

```
pd.Series(np.array([1, 2, 3]), index=['a', 'b', 'c'])  
# from a 1darray
```

PY

Try it Yourself

Or, we could use a dictionary to specify the index with keys:

```
pd.Series({'a': 1, 'b': 2, 'c':3}) # from a dict
```

PY

Try it Yourself

If we don't specify the index, by default, the index would be the integer positions starting from 0.

In a Series, we can access the value by its index directly:

```
series = pd.Series({'a': 1, 'b': 2, 'c':3})  
series['a']
```

PY

Try it Yourself

Accessing the value by its index, rather than the integer position comes in handy when the dataset is of thousands, if not millions, of rows. Series is the building block for the DataFrame we will introduce next.



Think of Series as numpy 1darray with index or row names.

## DataFrames

In data science, data is usually more than one-dimensional, and of different data types; thus Series is not sufficient. **DataFrames** are 2darrays with both row and column labels. One way to create a DataFrame from scratch is to pass in a **dict**. For example, this week, we sold 3 bottles of red wine to Adam, 6 to Bob, and 5 to Charles. We sold 5 bottles of white wine to Adam, 0 to Bob and 10 to Charles. We can organize the data into a DataFrame by creating a dict 'wine\_dict' with the number of bottles of each wine type we sold, then pass it along with the customer names as index to create a DataFrame 'sales'.

```
wine_dict = {  
    'red_wine': [3, 6, 5],  
    'white_wine':[5, 0, 10]  
}  
sales = pd.DataFrame(wine_dict, index=["adam", "bob",  
    "charles"])
```

PY

Try it Yourself

	red_wine	white_wine
adam	3	5
bob	6	0
charles	5	10

Think of DataFrame as a collection of the Series. Here, sales consists of two Series, one named under "red\_wine", the other "white\_wine", thus, we can access each series by calling its name:

```
sales['white_wine']
```

PY

Try it Yourself

We will see other ways to index into DataFrames in later parts.



If we don't supply index, the DataFrame will generate an integer index starting from 0.

## Inspect a DataFrame - Shape and Size

Let's take a look at a new DataFrame, in addition to heights and ages of the presidents, there is information on the order, names and parties. The DataFrame **presidents\_df** is read from a CSV file as follows. Note that index is set to be the names of presidents.

```
import pandas as pd
presidents_df =
pd.read_csv('https://sololearn.com/uploads/files/presid
ent_heights_party.csv', index_col='name')
```

PY

Similar to numpy, to get the dimensions of a DataFrame, use **.shape**

```
presidents_df.shape
```

PY

Try it Yourself

There are 45 rows and 4 columns in this DataFrame. To get the number of rows we can access the first element in the tuple.

```
presidents_df.shape[0]
```

PY

Try it Yourself

Size also works on DataFrame to return an integer representing the number of elements in this object.

```
presidents_df.size
```

PY

Try it Yourself



Here both methods. **.shape** and **.size**, work in the same way as with numpy ndarrays.

## Inspect a DataFrame - Head and Tail

Instead of looking at the entire dataset, we can just take a peep. To see the first few lines in a DataFrame, use `.head()`; if we don't specify `n` (the number of lines), by default, it displays the first five rows. Here we want to see the top 3 rows.

```
presidents_df.head(n=3)
```

py

Try it Yourself

	order	age	height	party
name				
George Washington	1	57	189	none
John Adams	2	61	170	federalist
Thomas Jefferson	3	57	189	democratic-republican

In `presidents_df`, the index is the name of the president, there are four columns: `order`, `age`, `height`, and `party`. Similarly, if we want to see the last few rows, we can use `.tail()`, the default is also five rows.

```
presidents_df.tail()
```

PY[Try it Yourself](#)

	order	age	height	party
name				
George H. W. Bush	41	64	188	republican
Bill Clinton	42	46	188	democratic
George W. Bush	43	54	182	republican
Barack Obama	44	47	185	democratic
Donald J. Trump	45	70	191	republican



Function `.head` is useful for quickly testing if the object has the right type of data in it.

## Inspect a DataFrame - Info

Use `.info()` to get an overview of the DataFrame. Its output includes index, column names, count of non-null values, dtypes, and memory usage.

```
presidents_df.info()
```

PY[Try it Yourself](#)

The dtype for order, age, and height is integers, while party is an object. The count of non-null values in each column is the same as the number of rows, indicating no missing values.



In addition to shape and size as seen in numpy, pandas features allow additional functionality to examine the data.

## Rows with .loc

Instead of memorizing the integer positions to locate the order, age, height, and party information of Abraham Lincoln, with DataFrame, we can access it by the name using **.loc**

```
presidents_df.loc['Abraham Lincoln']
```

PY

Try it Yourself

The result is a pandas Series of shape (4,).

```
type(presidents_df.loc['Abraham Lincoln'])  
presidents_df.loc['Abraham Lincoln'].shape
```

PY

Try it Yourself

We can also slice by index. Say we are interested in gathering information on all of the presidents between Abraham Lincoln and Ulysses S. Grant:

```
presidents_df.loc['Abraham Lincoln':'Ulysses S. Grant']
```

PY

Try it Yourself

	order	age	height	party
name				
<b>Abraham Lincoln</b>	16	52	193	republican
<b>Andrew Johnson</b>	17	56	178	national union
<b>Ulysses S. Grant</b>	18	46	173	republican

The result is a new DataFrame, a subset of 'presidents\_df'.



.loc[ ] allows us to select data by label or by a conditional statement.



## Rows with .iloc

Alternatively, if we do know the integer position(s), we can use `.iloc` to access the row(s).

```
presidents_df.iloc[15]
```

PY

Try it Yourself

To gather information from the 16th to 18th presidents, we can then:

```
presidents_df.iloc[15:18]
```

PY

Try it Yourself

	order	age	height	party
name				
Abraham Lincoln	16	52	193	republican
Andrew Johnson	17	56	178	national union
Ulysses S. Grant	18	46	173	republican



Both `.loc[ ]` and `.iloc[ ]` may be used with a boolean array to subset the data.

## Columns

We can retrieve an entire column from `presidents_df` by name. First we access all the column names:

```
presidents_df.columns
```

PY

Try it Yourself

Which returns an index object containing all column names. Then we can access the column height by:

```
presidents_df['height']  
presidents_df['height'].shape
```

PY

Try it Yourself

Which returns a Series containing heights from all U.S. presidents.

To select multiple columns, we pass the names in a list, resulting in a DataFrame. Remember, we can use `.head()` to access the first 3 rows as shown below:

```
presidents_df[['height', 'age']].head(n=3)
```

PY

Try it Yourself

	height	age
name		
George Washington	189	57
John Adams	170	61
Thomas Jefferson	189	57

In such a way, we focus on only the columns of interest.



When accessing a single column, one bracket results in a Series (single dimension) and double brackets results in a DataFrame (multi dimensional).

## More with .loc

If we wanted to access columns order, age, and height, we can do it with `.loc`. `.loc` allows us to access any of the columns. For example, if we wanted to access columns from order through height for the first three presidents:

```
presidents_df.loc[:, 'order':'height'].head(n=3)
```

PY

Try it Yourself

	order	age	height
name			
George Washington	1	57	189
John Adams	2	61	170
Thomas Jefferson	3	57	189



The index in pandas makes retrieving information from rows or columns convenient and easy, especially when the data set is large or there are many columns. Therefore, we don't have to memorize the integer positions of each row or column.

## Min / Max / Mean

It's not practical to print out an entire dataset with a large sample size. Instead, we want to summarize and characterize sample data using only a few values. Summary statistics include **measures of location** and **measures of spread**. Measures of location are quantities that represent the average value of a variable while measures of spread represent how similar or dissimilar the values of a variable are.

**Measures of Location** - Minimum, Maximum, Mean

**Measures of Spread** - Range, Variance, Standard Deviation

The simplest summary statistics, which are measures of location, include the minimum, the smallest number:

```
presidents_df.min()
```

PY

Try it Yourself

maximum, the largest number:

```
presidents_df.max()
```

PY

Try it Yourself

and mean, the average:

```
presidents_df.mean()
```

PY

Try it Yourself

Recall the arithmetic mean is the sum of the elements divided by the number of elements, in python 3.x, division of integers results in a float number.

Once the minimum and maximum are known, we can determine the **range**, a measure of spread. For example, the height for all U.S. presidents ranges from 163 -- 193 cm.

The mean tells us where the data is centered. For instance, the average age at the start of the presidency is 54.71 years. Note that `mean()` can only operate on the numeric values, thus the column 'party' was omitted.



These methods work on Series as well. For example, 'presidents\_df['age'].mean()' also results in 54.71

## Quantiles

**Quantiles** are cut points dividing the range of the data into continuous intervals with an equal number of observations. Median is the only cut point in 2-quantiles, such that 50% of the data is below the median with the other half above it.

**Quantiles** let us quickly divide a set of data into four groups, making it easy to see which of the four groups a particular data point is in. Quartiles are then 4-quantiles, that is, 25% of the data are between the minimum and first quartile, the next is 25% between the first quartile and median, the next 25% is between the median and the third quartile, and the last 25% of the data lies between the third quartile and the maximum.

```
presidents_df['age'].quantile([0.25, 0.5, 0.75, 1])
```

PY

Try it Yourself

Here 25% of presidents started their presidency at 51 years old or younger, while half started their presidency at 55 years old or younger.

**Mean** and **median** are usually not of the same value, unless the data is perfectly symmetric. The mean is the average of all the numbers added together and divided by the amount of numbers added. The median is the value separating the higher half from the lower half of the data sample. In the age data, the mean is close to its median, this implies that the data might be symmetric.

```
presidents_df['age'].mean()  
presidents_df['age'].median()
```

PY

Try it Yourself



Both `.quantile(0.5)` and `.median()` result in the same output.

## Variance and Standard Deviation

In probability and statistics, **variance** is the mean squared deviation of each data point from the mean of the entire dataset.

You can think of it as how far apart a set of numbers are spread out from their average value. **Standard deviation** (std) is the square root of variance. A high std implies a large spread, and a low std indicates a small spread, or most points are close to the mean.

In one extreme example, the data consists of all constant 2, there is no variation, thus the variation is 0.0, so is its std:

```
const = pd.Series([2, 2, 2])  
const.var()  
const.std()
```

PY

Try it Yourself

Lets consider another example:

```
[2, 3, 4]
```

The mean of [2,3,4] is  $(2+3+4)/3 = 3.0$ , and its variation is  $(2-3)^2 + (3-3)^2 + (4-3)^2 = 1+0+1 = 2$ . Note that in Python, `.var()` will return the variance divided by  $N-1$  where  $N$  is the length of the data, the output is then  $2/(3-1) = 1$ .

```
dat = pd.Series([2, 3, 4])
dat.mean()
dat.var()
# Outputs:
# 3.0
# 1.0
```

PY

Try it Yourself

And the std is just the square root of variance:

```
dat.std()
# Outputs: 1.0
```

PY

Try it Yourself



For the ages of the presidents:

```
presidents_df['age'].var()  
presidents_df['age'].std()  
# Outputs:  
# 43.5  
# 6.595354
```

PY

Try it Yourself

We can apply std on the entire DataFrame to get column-wise standard deviation.

```
presidents_df.std()  
# Outputs:  
# order: 13.136502  
# age: 6.595354  
# height: 6.977236
```

PY

Try it Yourself



We can apply min, max, quantile, and var on the entire DataFrame in the same way.

## describe()

`describe()` prints out almost all of the **summary statistics** mentioned previously except for the variance. In addition, it counts all non-null values of each column.

```
presidents_df['age'].describe()  
presidents_df.describe()
```

PY

Try it Yourself

	order	age	height
count	45.000000	45.000000	45.000000
mean	23.022222	55.000000	180.000000
std	13.136502	6.595453	6.977236
min	1.000000	42.000000	163.000000
25%	12.000000	51.000000	175.000000
50%	23.000000	55.000000	182.000000
75%	34.000000	58.000000	183.000000
max	45.000000	70.000000	193.000000

From the output we can see that there are 45 non-null data points of ages, with a **mean** 55 and **std** 6.60. The ages **range** from 42 to 70 with a **median** 55. Its first and third **quartiles** are 51 and 58, respectively. Now we have an overall description of all age data. In addition to being applied to a series, `describe()` can be applied to a DataFrame with multiple columns.

As the count (=45) suggests, there are no null values in any of the three columns. Order is simply an index from 1 to 45. Interestingly, both age and height lie in the interval of roughly the same length,  $70-42 = 28$  for age while  $193-163 = 30$  for height. Also both features are of similar standard deviations, indicating a similar spread of the data.



**.describe()** ignores the null values, such as `NaN` (Not a Number) and generates the descriptive statistics that summarize the central tendency (i.e., mean), dispersion (i.e., standard deviation), and shape (i.e., min, max, and quartiles) of a dataset's distribution.

## Categorical Variable

The fourth column 'party' was omitted in the output of `.describe()` because it is a **categorical variable**. A categorical variable is one that takes on a single value from a limited set of categories. It doesn't make sense to calculate the mean of democratic, republican, federalist, and other parties. We can check the unique values and corresponding frequency by using `.value_counts()`:

```
presidents_df['party'].value_counts()
```

PY

Try it Yourself

We can also call `.describe()` to see that there are 45 non-null values, 7 unique parties, the most frequent party is republican, with a total of 19 presidents belonging to this party.

```
presidents_df['party'].describe()
```

PY

Try it Yourself



Summary statistics provides us with a large amount of information put as simply as possible. The measure of location, median, is more robust than mean, for continuous variables as the latter is sensitive to outliers, e.g., extremely large values.

## Groupby

Summary statistics on an entire dataset provides a good overall view, but often we're interested in some calculation conditional upon a given label or category. For example, what is the average height conditional of the presidents party?

To find the value based on a condition, we can use the **groupby** operation. Think of groupby doing three steps: split, apply, and combine. The split step breaks the DataFrame into multiple DataFrames based on the value of the specified key; the apply step is to perform the operation inside each smaller DataFrame; the last step combines the pieces back into the larger DataFrame.

```
presidents_df.groupby('party')
```

PY

The `.groupby("party")` returns a `DataFrameGroupBy` object, not a set of DataFrames. To produce a result, apply an aggregate (`.mean()`) to this `DataFrameGroupBy` object:

```
presidents_df.groupby('party').mean()
```

PY

Try it Yourself

	order	age	height
party			
democratic	26.066667	52.600000	181.066667
democratic-republican	4.500000	57.250000	176.500000
federalist	2.000000	61.000000	170.000000
national union	17.000000	56.000000	178.000000
none	1.000000	57.000000	189.000000
republican	29.631579	55.263158	180.894737
whig	11.000000	58.250000	176.000000



The `mean()` method is one of many possibilities, you can apply any pandas or numpy aggregation function, or any DataFrame operation, as we demonstrate through this course.

## Aggregation

We can also perform multiple operations on the groupby object using `.agg()` method. It takes a string, a function, or a list thereof. For example, we would like to obtain the min, median, and max values of heights grouped by party:

```
presidents_df.groupby('party')['height'].agg(['min',  
np.median, max])
```

PY

Try it Yourself

	min	median	max
party			
democratic	168	180	193
democratic-republican	163	177	189
federalist	170	170	170
national union	178	178	178
none	189	189	189
republican	168	182	193
whig	173	174	183

From the output we can see, the heights of the democratic presidents range from 168 cm to 193 cm, with a median at 180 cm.

Often time we are interested in different summary statistics for multiple columns. For instance, we would like to check the median and mean of heights, but minimum and maximum for ages, grouped by party. In this case, we can pass a dict with key indicate the column name, and value indicate the functions:

```
presidents_df.groupby('party')  
    .agg({'height': [np.median, np.mean],  
         'age': [min, max]})
```

PY

Try it Yourself

	height		age	
	median	mean	min	max
party				
democratic	180	181.066667	43	65
democratic-republican	177	176.500000	57	58
federalist	170	170.000000	61	61
national union	178	178.000000	56	56
none	189	189.000000	57	57
republican	182	180.894737	42	70
whig	174	176.000000	50	68



Using **groupby** and **agg** provides us the flexibility and therefore the power to look into various perspectives of a variable or column conditioned on categories.